

A Comparative Analysis of Machine Learning Models for Predicting Student Performance

Reine Nassour
Department of Computer Science
Lebanese American University
reinenassour341@gmail.com

Abstract

Student academic performance prediction has become an important topic in educational data mining and artificial intelligence. Early identification of students who may experience academic difficulties can help educators implement appropriate interventions and improve learning outcomes. In recent years, machine learning techniques have been widely used to analyze educational datasets and predict student performance based on various academic and behavioral factors.

This study investigates the effectiveness of multiple machine learning algorithms for predicting student academic performance. Several regression models, including Linear Regression, Decision Trees, Random Forest, K-Nearest Neighbors, Gradient Boosting, Extra Trees, XGBoost, and Support Vector Regression, were implemented and compared. The models were trained using a student performance dataset containing several features related to academic and contextual factors.

To evaluate the predictive performance of the models, the coefficient of determination (R^2) was used as the primary evaluation metric. The experimental results provide a comparative analysis of the algorithms and highlight the models that achieve the highest predictive accuracy. The findings demonstrate the potential of machine learning techniques in supporting data-driven decision-making in the educational domain.

Keywords

Machine Learning, Student Performance Prediction, Educational Data Mining, Regression Models, Artificial Intelligence

1. Introduction

Predicting student academic performance has become a significant research topic in the field of educational data mining and learning analytics. Educational institutions continuously seek effective methods to monitor student progress and identify students who may face academic challenges. Accurate prediction of student performance can enable educators to provide early interventions, improve learning strategies, and enhance overall academic outcomes. With the rapid growth of educational data generated from learning systems, machine learning techniques have emerged as powerful tools for analyzing and extracting valuable insights from such data.

In recent years, machine learning algorithms have been widely applied to educational datasets to predict student success or failure based on various academic, demographic, and behavioral features. These techniques allow researchers to build predictive models capable of identifying patterns and relationships that may not be easily detected using traditional statistical methods. As a result, machine learning has become an essential approach in supporting data-driven decision-making in modern educational environments.

Several studies have explored the use of different machine learning models for predicting student performance. Algorithms such as Decision Trees, Random Forest, Support Vector Machines, and K-Nearest Neighbors have demonstrated promising results in educational prediction tasks. However, the performance of these algorithms may vary depending on the dataset characteristics, feature selection, and model configuration. Therefore, conducting comparative studies among multiple algorithms is important to determine which models perform best for a given problem.

This study aims to develop and compare multiple machine learning algorithms for predicting student academic performance. A variety of regression models, including Linear Regression, Random Forest, K-Nearest Neighbors, Decision Trees, Gradient Boosting, Extra Trees, XGBoost, and Support Vector Regression, were implemented and evaluated. The performance of these models was assessed using the coefficient of determination (R^2). The main objective of this research is to identify the most effective model for predicting student performance and to provide insights into the applicability of machine learning techniques in the educational domain.

2. Related Work

In recent years, educational data mining and learning analytics have gained increasing attention as researchers attempt to improve educational outcomes using data-driven approaches. Predicting student academic performance has become one of the most widely studied topics in this field. Machine learning techniques have been extensively applied to analyze educational datasets and identify factors that influence student success or failure.

Several studies have explored the use of machine learning algorithms for predicting student performance. For example, Decision Tree and Random Forest models have been widely used due to their ability to capture complex relationships between input variables and target outcomes. These models are particularly effective when dealing with structured educational data and can provide insights into the importance of different features affecting student performance.

Support Vector Machines and K-Nearest Neighbors have also been applied in many studies related to student performance prediction. These algorithms are capable of handling nonlinear relationships between variables and have demonstrated promising results in classification and regression tasks. Additionally, ensemble learning methods such as Gradient Boosting and Extra Trees have been shown to improve predictive accuracy by combining multiple decision trees into a single model.

More recently, advanced machine learning models such as Extreme Gradient Boosting (XGBoost) have become increasingly popular due to their high predictive performance and computational efficiency. These models have been successfully applied in various prediction tasks, including educational data analysis, where they can effectively model complex relationships within datasets.

Despite the large number of studies in this area, the performance of machine learning models may vary depending on the dataset characteristics, feature engineering techniques, and evaluation methods used. Therefore, comparative studies that evaluate multiple machine learning algorithms on the same dataset are essential to determine which models provide the most accurate predictions. This research contributes to this area by implementing and comparing a diverse set of regression algorithms to predict student academic performance.

3. Dataset Description

The dataset used in this study was obtained from the Kaggle platform, which provides publicly available datasets for machine learning research.

The dataset used in this study focuses on predicting student academic performance using several educational and demographic factors. The dataset contains **708 student records** and **10 variables**, representing different characteristics that may influence academic outcomes.

Each record corresponds to an individual student and includes both numerical and categorical attributes describing study behavior, learning conditions, and previous academic performance. The dataset includes features such as weekly study hours, attendance rate, past examination scores, parental education level, internet access at home, and participation in extracurricular activities.

The primary target variable in this study is the **Final_Exam_Score**, which represents the final academic performance of each student. The machine learning models aim to predict this score based on the available input features.

Several features in the dataset provide valuable insights into factors that may influence student performance. For example, the variable **Study_Hours_per_Week** reflects the amount of time a student dedicates to studying, while **Attendance_Rate** indicates the level of student participation in classes. Similarly, **Past_Exam_Scores** represent previous academic achievement, which may strongly influence future performance.

In addition to academic indicators, the dataset also includes contextual variables such as **Parental_Education_Level**, **Internet_Access_at_Home**, and **Extracurricular_Activities**, which may indirectly affect student learning outcomes. These variables allow machine learning models to capture complex relationships between academic performance and environmental factors.

Before training the machine learning models, the dataset was carefully examined and prepared through several preprocessing steps to ensure data quality and model compatibility.

Feature	Description
Study_Hours_per_Week	Number of hours the student studies per week
Attendance_Rate	Percentage of class attendance
Past_Exam_Scores	Scores obtained in previous exams
Parental_Education_Level	Education level of the student's parents
Internet_Access_at_Home	Whether the student has internet access at home

Extracurricular_Activities	Participation in activities outside the classroom
Final_Exam_Score	Final academic score obtained by the student

4. Methodology

This section describes the dataset used in this study, the preprocessing steps applied to the data, the machine learning models implemented for prediction, and the evaluation metrics used to assess model performance.

4.1 Dataset

The dataset used in this study focuses on predicting student academic performance based on several educational and contextual factors. The dataset contains multiple observations representing individual students, along with several features that describe characteristics related to academic activities and learning conditions.

These features may include variables such as study time, attendance, previous academic performance, and other factors that may influence student outcomes. The target variable represents the final student performance score, which the machine learning models attempt to predict.

The dataset was divided into two subsets: a training set and a testing set. The training dataset was used to train the machine learning models, while the testing dataset was used to evaluate the predictive performance of the trained models on unseen data.

4.2 Data Preprocessing

Before training the machine learning models, several preprocessing steps were applied to prepare the data for analysis. Data preprocessing is an essential step in machine learning as it ensures that the dataset is clean, consistent, and suitable for model training.

First, the dataset was examined to identify any missing values or inconsistencies. Appropriate techniques were applied to handle these issues when necessary. Next, feature scaling was applied to standardize the values of the input variables. This step is particularly important for algorithms such as K-Nearest Neighbors and Support Vector Regression, which are sensitive to differences in feature scales.

The dataset was then split into training and testing sets to ensure that model evaluation is performed on data that was not used during the training phase.

4.3 Machine Learning Models

To evaluate the effectiveness of machine learning techniques for predicting student performance, multiple regression algorithms were implemented and compared in this study.

The models used in this research include:

- Linear Regression
- Random Forest Regressor
- K-Nearest Neighbors Regressor
- Radius Neighbors Regressor
- Decision Tree Regressor
- Extra Tree Regressor
- Extra Trees Regressor
- Gradient Boosting Regressor
- Extreme Gradient Boosting (XGBoost)
- Support Vector Regression (SVR) with different kernels
- Nu-Support Vector Regression (NuSVR)

These models were selected because they represent different machine learning approaches, including linear models, tree-based methods, ensemble learning techniques, and kernel-based models. By comparing these algorithms, it becomes possible to determine which approach provides the most accurate predictions for the given dataset.

4.4 Model Evaluation

To evaluate the predictive performance of the implemented machine learning models, the **coefficient of determination (R^2)** was used as the primary evaluation metric. This metric measures how well the predicted values produced by the model match the actual observed values.

The coefficient of determination is defined as:

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \bar{y})^2}{\sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

where y_i represents the actual values, $\hat{y}_i$ represents the predicted values, and $\bar{y}$ represents the mean of the observed values. An R^2 value closer to 1 indicates better predictive performance, meaning that the model explains a larger proportion of the variance in the target variable.

5. Experiments and Results

5.1 Experimental Setup

In order to evaluate the performance of the machine learning models, several regression algorithms were implemented using the Python programming language and the Scikit-learn machine learning library. The models were trained using the training dataset and evaluated using the testing dataset to assess their ability to generalize to unseen data.

Each algorithm was trained using the same dataset and preprocessing procedures to ensure a fair comparison between the models. The predictive performance of each model was evaluated using the coefficient of determination (R^2), which measures how well the predicted values correspond to the actual target values.

5.2 Model Performance Comparison

After training the models, the predictive performance of each algorithm was calculated using the R^2 metric. Table 1 presents the performance of the implemented models.

Model	R^2 Score
Linear Regression	0.7769
Random Forest	0.7125
K-Nearest Neighbors	0.6845
Radius Neighbors	0.5153
Decision Tree	0.4213
Extra Tree	0.4706
Extra Trees	0.7184
Gradient Boosting	0.6901
XGBoost	0.6278
SVR (RBF Kernel)	0.7619
SVR (Linear Kernel)	0.7759
SVR (Polynomial Kernel)	0.7434
NuSVR	0.7396

The results indicate that different machine learning algorithms exhibit varying levels of predictive performance. Ensemble learning methods such as Random Forest, Gradient Boosting, and XGBoost typically demonstrate strong predictive capabilities due to their ability to combine multiple decision trees and capture complex patterns within the data.

5.3 Visualization of Model Performance

To provide a clearer comparison between the machine learning models, a graphical visualization of the R^2 scores was generated. The comparison graph illustrates the predictive performance of each algorithm.




Figure 1. Comparison of machine learning algorithms based on R^2 performance.

To provide a clearer comparison between the machine learning models, a graphical visualization of the R^2 scores was generated. **As shown in Figure 1**, the comparison graph illustrates the predictive performance of each algorithm.




Figure 2. Actual vs Predicted Student Performance.

This scatter plot compares the actual student exam scores with the predicted values generated by the machine learning model. A closer alignment of points along the diagonal trend indicates better predictive performance.

5.4 Discussion of Results

Based on the experimental results, it can be observed that some machine learning algorithms outperform others in predicting student academic performance. Ensemble methods such as Random Forest, Gradient Boosting, and XGBoost often achieve higher accuracy because they combine multiple models and reduce prediction errors.

On the other hand, simpler models such as Linear Regression may not capture complex relationships between variables as effectively as ensemble or nonlinear models. Additionally, algorithms such as Support Vector Regression and K-Nearest Neighbors can perform well depending on the characteristics of the dataset and the chosen hyperparameters.

Overall, the results demonstrate that machine learning techniques provide an effective approach for predicting student academic performance and can support educational institutions in identifying students who may require additional academic support.

7. Conclusion

This study investigated the application of machine learning techniques for predicting student academic performance using a dataset containing various academic and contextual features. Multiple regression algorithms were implemented and compared, including Linear Regression, Random Forest, K-Nearest Neighbors, Decision Trees, Extra Trees, Gradient Boosting, XGBoost, and Support Vector Regression.

The experimental results demonstrated that machine learning models can effectively capture relationships between educational factors and student performance. The comparison of different algorithms revealed variations in predictive accuracy, highlighting the importance of selecting appropriate models for educational prediction tasks. Ensemble-based methods such as Random Forest, Gradient Boosting, and XGBoost generally showed strong predictive capabilities due to their ability to model complex patterns in the data.

The findings of this research confirm the potential of machine learning approaches in supporting educational institutions and educators in identifying students who may require additional academic support. Early prediction of student performance can contribute to improving learning outcomes and enabling data-driven decision-making in educational environments.

Future research may explore the integration of additional datasets, advanced feature engineering techniques, and deep learning approaches to further improve prediction accuracy. Additionally, incorporating real-time educational data from learning management systems could provide deeper insights into student learning behaviors and performance patterns.

To ensure reproducibility and transparency of the experimental results, the implementation of the machine learning models and the complete experimental setup used in this study are publicly available in the following repository:

Code Repository:

<https://github.com/ReineNassour/Machine-Learning-Models-for-Predicting-Student-Academic-Performance>

References

- Amr Maree. (2024). *Student Performance Prediction Dataset*. Kaggle.
<https://www.kaggle.com/datasets/amrmaree/student-performance-prediction>
- Breiman, L. (2001). *Random Forests*. Machine Learning.
- Chen, T., & Guestrin, C. (2016). *XGBoost: A scalable tree boosting system*.
- Cortes, C., & Vapnik, V. (1995). *Support Vector Networks*.